

```

# This Python 3 environment comes with many helpful analytics
libraries installed
# It is defined by the kaggle/python Docker image:
https://github.com/kaggle/docker-python
# For example, here's several helpful packages to load

import numpy as np # linear algebra
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns# data processing, CSV file I/O (e.g.
pd.read_csv)

# Input data files are available in the read-only "../input/"
directory
# For example, running this (by clicking run or pressing Shift+Enter)
will list all files under the input directory

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

# You can write up to 20GB to the current directory (/kaggle/working/)
that gets preserved as output when you create a version using "Save &
Run All"
# You can also write temporary files to /kaggle/temp/, but they won't
be saved outside of the current session

/kaggle/input/taxi-fare-guru-total-amount-prediction-challenge/
sample.csv.csv
/kaggle/input/taxi-fare-guru-total-amount-prediction-challenge/train.c
sv
/kaggle/input/taxi-fare-guru-total-amount-prediction-challenge/test.cs
v

```

## 1. Data Loading

```

data=pd.read_csv('/kaggle/input/taxi-fare-guru-total-amount-
prediction-challenge/train.csv')
data_t=pd.read_csv('/kaggle/input/taxi-fare-guru-total-amount-
prediction-challenge/test.csv')

```

```
data.head()
```

```

  VendorID tpep_pickup_datetime tpep_dropoff_datetime
passenger_count \
0          1  2023-06-28 17:20:21    2023-06-28 16:34:45
1.0
1          0  2023-06-29 23:05:01    2023-06-29 22:01:35
1.0

```

2	1	2023-06-30 10:19:31	2023-06-30 11:13:10
1.0			
3	0	2023-06-29 13:23:09	2023-06-29 14:20:01
1.0			
4	1	2023-06-29 22:03:32	2023-06-29 22:22:22
3.0			

	trip_distance	RatecodeID	store_and_fwd_flag	PULocationID
DOLocationID \				
0	2.14	1.0	N	120
9				
1	2.70	1.0	N	15
215				
2	1.15	1.0	N	167
223				
3	0.40	1.0	N	128
239				
4	1.10	1.0	N	203
52				

	payment_type	extra	tip_amount	tolls_amount	improvement_surcharge
\					
0	Credit Card	2.5	7.165589	0.0	1.0
1	Credit Card	3.5	6.067401	0.0	1.0
2	Credit Card	0.0	4.111547	0.0	1.0
3	Credit Card	2.5	6.411079	0.0	1.0
4	Credit Card	1.0	4.769377	0.0	1.0

	total_amount	congestion_surcharge	Airport_fee
0	20.64	2.5	0.0
1	25.55	2.5	0.0
2	17.64	2.5	0.0
3	12.80	2.5	0.0
4	18.00	2.5	0.0

```
feature_list = data.columns[:-1].values
label = data.columns[-3]
print("Feature List:", feature_list)
print("Label:", label)
```

```
Feature List: ['VendorID' 'tpep_pickup_datetime'
'tpep_dropoff_datetime'
'passenger_count' 'trip_distance' 'RatecodeID' 'store_and_fwd_flag'
'PULocationID' 'DOLocationID' 'payment_type' 'extra' 'tip_amount'
'tolls_amount' 'improvement_surcharge' 'total_amount']
```

```
'congestion_surcharge']
Label: total_amount
```

## 2. Exploratory Data Analysis

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
# Let us summarize our data
data.describe()
```

	VendorID	passenger_count	trip_distance	RatecodeID	\
count	175000.000000	168923.000000	175000.000000	168923.000000	
mean	0.728377	1.357678	5.145930	1.518307	
std	0.445606	0.891283	394.971052	6.514678	
min	0.000000	0.000000	0.000000	1.000000	
25%	0.000000	1.000000	1.080000	1.000000	
50%	1.000000	1.000000	1.840000	1.000000	
75%	1.000000	1.000000	3.610000	1.000000	
max	2.000000	9.000000	135182.060000	99.000000	

	PULocationID	DOLocationID	extra	tip_amount	\
count	175000.000000	175000.000000	175000.000000	175000.000000	
mean	132.710349	132.701429	1.932143	6.127497	
std	76.148799	76.192493	1.948497	4.610834	
min	1.000000	1.000000	-7.500000	0.000079	
25%	67.000000	67.000000	0.000000	3.473321	
50%	133.000000	133.000000	1.000000	5.286217	
75%	199.000000	199.000000	2.500000	7.502746	
max	264.000000	264.000000	11.750000	484.876151	

	tolls_amount	improvement_surcharge	total_amount	\
count	175000.000000	175000.000000	175000.000000	
mean	0.646816	0.979689	29.633901	
std	2.328274	0.198775	25.425206	
min	-29.300000	-1.000000	-576.750000	
25%	0.000000	1.000000	16.300000	
50%	0.000000	1.000000	21.450000	
75%	0.000000	1.000000	31.800000	
max	80.000000	1.000000	587.250000	

	congestion_surcharge	Airport_fee
count	168923.000000	168923.000000
mean	2.246971	0.158825
std	0.819216	0.511968
min	-2.500000	-1.750000
25%	2.500000	0.000000
50%	2.500000	0.000000

75%	2.500000	0.000000
max	2.500000	1.750000

*# Finding out the missing values in our data*

```
missing_value_counts = data.isnull().sum()
missing_value_counts
```

```
VendorID      0
tpep_pickup_datetime  0
tpep_dropoff_datetime  0
passenger_count  6077
trip_distance    0
RatecodeID    6077
store_and_fwd_flag  6077
PULocationID    0
DOLocationID    0
payment_type    0
extra           0
tip_amount      0
tolls_amount    0
improvement_surcharge  0
total_amount    0
congestion_surcharge  6077
Airport_fee     6077
dtype: int64
```

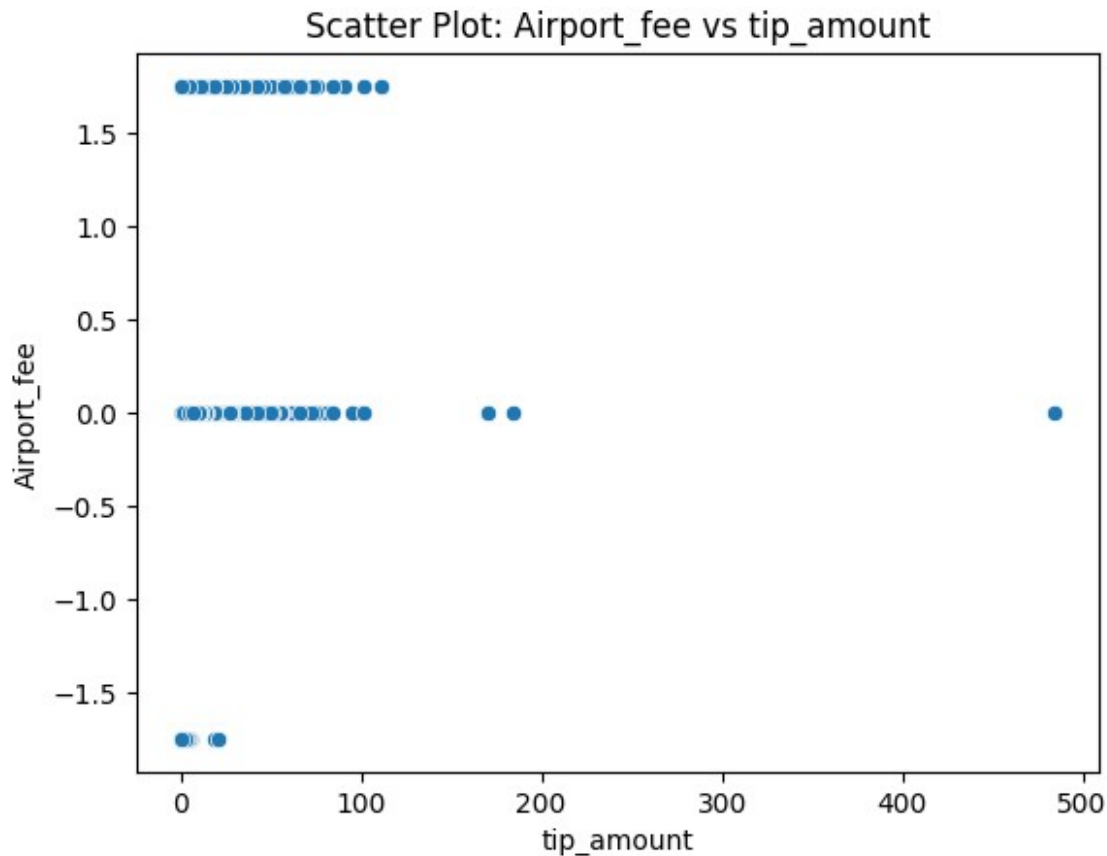
*# counts the occurrences of unique values in the Airport\_fee column.*

```
unique_values = data['Airport_fee'].value_counts()
unique_values
```

```
Airport_fee
0.00      153074
1.75       15590
-1.75        259
Name: count, dtype: int64
```

*# Scatter plot between airport fee and tip\_amount*

```
sns.scatterplot(x='tip_amount', y='Airport_fee', data=data)
plt.title('Scatter Plot: Airport_fee vs tip_amount')
plt.xlabel('tip_amount')
plt.ylabel('Airport_fee')
plt.show()
```




*The scatter plot effectively illustrates the relationship between 'tip\_amount' and 'Airport\_fee'. It allows us to identify outliers within the dataset, as certain data points deviate significantly from the overall pattern. Besides, it can also be concluded that airport\_fee is a categorical data.*

```
# Using heatmap to find correlation among the specified features
selected_features = data[['passenger_count', 'trip_distance', 'extra',
'total_amount', 'improvement_surcharge', 'tip_amount', 'tolls_amount']]

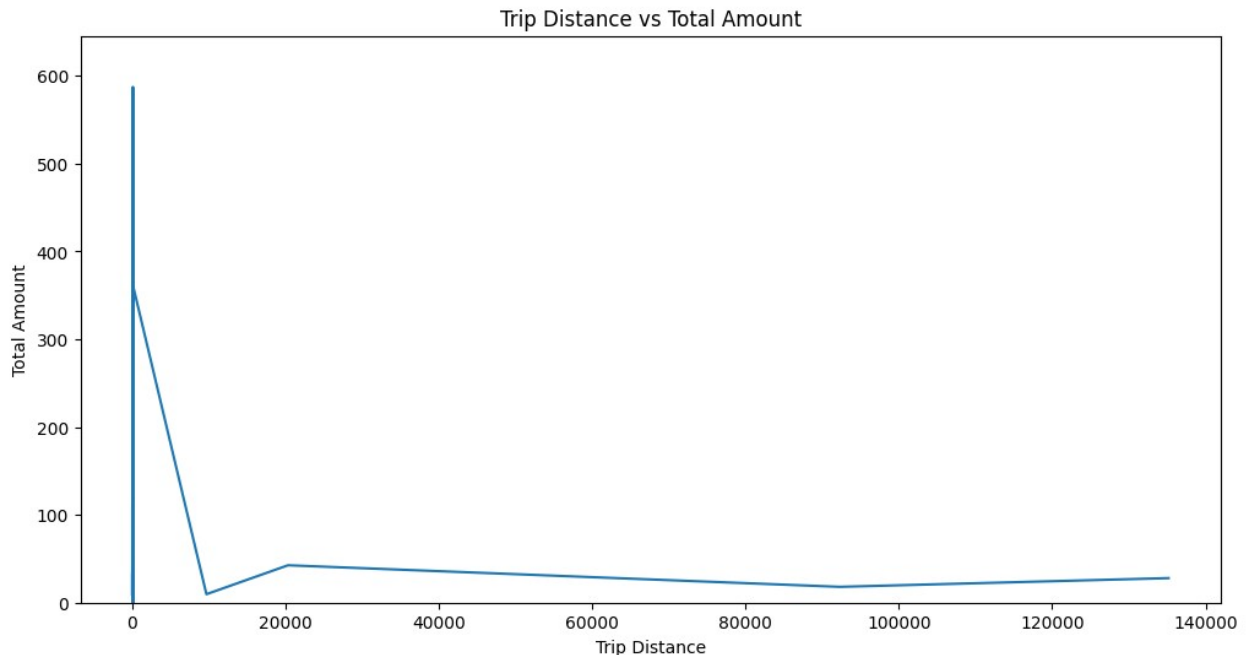
correlation_matrix = selected_features.corr()

plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm',
fmt=".2f", linewidths=.5)
plt.title('Correlation Matrix of Selected Features')
plt.show()
```



*In our analysis, it is evident that tolls amount and total amount exhibit a high degree of correlation, suggesting a strong relationship between these two variables. On the other hand, features such as passenger count and extras appear to be the least correlated, indicating a weaker association between these aspects.*

```
plt.figure(figsize=(12, 6))
sns.lineplot(x='trip_distance', y='total_amount', data=data)
plt.title('Trip Distance vs Total Amount')
plt.xlabel('Trip Distance')
plt.ylabel('Total Amount')
plt.ylim(0, None)
plt.show()
```



*As the distance increases, a discernible trend emerges, indicating an unchanging pattern in the total trip amount. This suggests that, beyond a certain distance, the total amount does not exhibit significant variation, implying a potential saturation or consistency in pricing with increasing trip distance.*


```
# Let us calculate the percentage each unique payment mode present
over there in the dataset.
```

```
payment_method_c = data['payment_type'].value_counts()
payment_method_p = (payment_method_c / len(data)) * 100
payment_method_p_rounded = payment_method_p.round(2)
print(payment_method_p_rounded)
```

```
payment_type
Credit Card    77.29
Cash           17.22
Wallet          3.47
unknown        1.33
UPI             0.68
Name: count, dtype: float64
```

```
# Here is our plot for percentage distribution of the various modes of
payment
```

```
plt.figure(figsize=(10, 6))
sns.barplot(x=payment_method_p_rounded.index,
y=payment_method_p_rounded.values)
plt.title('Percentage Distribution of Payment Types')
plt.xlabel('Payment Type')
plt.ylabel('Percentage')
plt.show()
```



*Clearly, the number of people using credit card is more when compared to the other modes like upi and wallet*

```
''' Converting the 'tpep_pickup_datetime' and 'tpep_dropoff_datetime'
columns to datetime objects
and then extracting day, month, and hour information for both pickup
and dropoff times. '''
```

```
data['tpep_pickup_datetime'] =
pd.to_datetime(data['tpep_pickup_datetime'])
data['tpep_dropoff_datetime'] =
pd.to_datetime(data['tpep_dropoff_datetime'])
```

```
data['pickup_day'] = data['tpep_pickup_datetime'].dt.day
data['pickup_month'] = data['tpep_pickup_datetime'].dt.month
data['pickup_hour'] = data['tpep_pickup_datetime'].dt.hour

data['dropoff_day'] = data['tpep_dropoff_datetime'].dt.day
data['dropoff_month'] = data['tpep_dropoff_datetime'].dt.month
data['dropoff_hour'] = data['tpep_dropoff_datetime'].dt.hour
```

```
data.drop(columns=['tpep_pickup_datetime', 'tpep_dropoff_datetime'],
inplace=True)
```



```

data_t['tpep_pickup_datetime'] =
pd.to_datetime(data_t['tpep_pickup_datetime'])
data_t['tpep_dropoff_datetime'] =
pd.to_datetime(data_t['tpep_dropoff_datetime'])

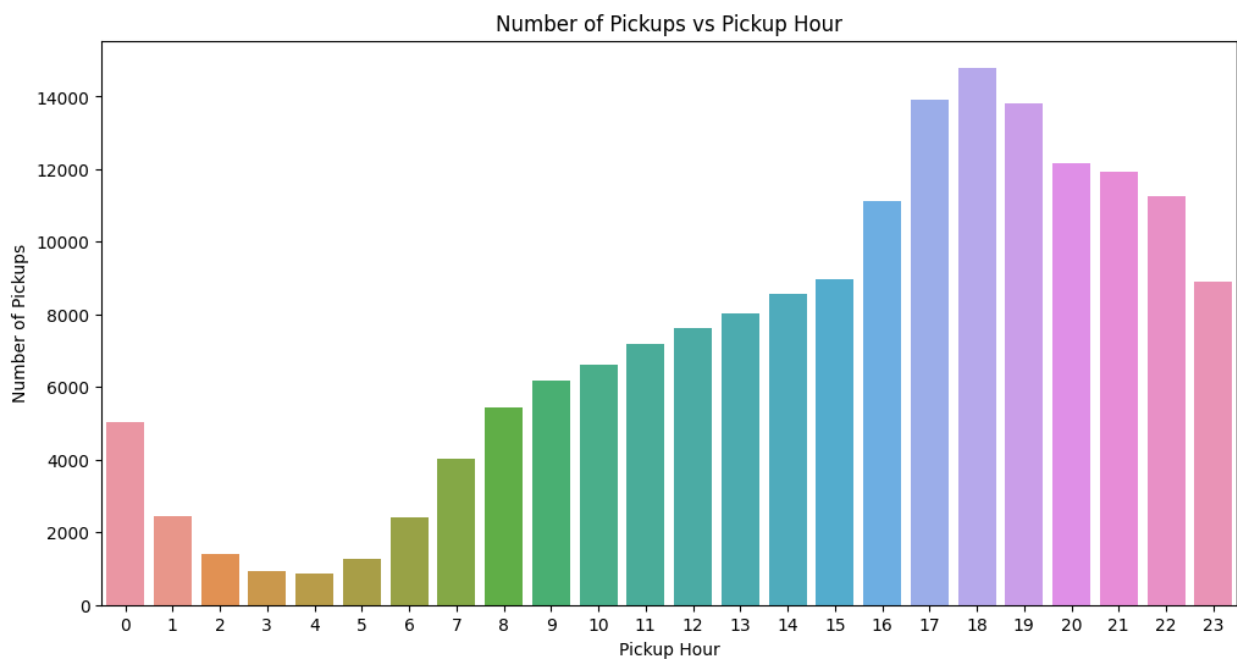
data_t['pickup_day'] = data_t['tpep_pickup_datetime'].dt.day
data_t['pickup_month'] = data_t['tpep_pickup_datetime'].dt.month
data_t['pickup_hour'] = data_t['tpep_pickup_datetime'].dt.hour

data_t['dropoff_day'] = data_t['tpep_dropoff_datetime'].dt.day
data_t['dropoff_month'] = data_t['tpep_dropoff_datetime'].dt.month
data_t['dropoff_hour'] = data_t['tpep_dropoff_datetime'].dt.hour

data_t.drop(columns=['tpep_pickup_datetime', 'tpep_dropoff_datetime'],
inplace=True)

# Here is a plot for number of pickups vs pickup hour
plt.figure(figsize=(12, 6))
sns.countplot(x='pickup_hour', data=data)
plt.title('Number of Pickups vs Pickup Hour')
plt.xlabel('Pickup Hour')
plt.ylabel('Number of Pickups')
plt.show()

```



*When it comes to pickup hour, 2 am to 6 am in the morning is quiet idle where 5 pm to 10 pm in the evening seems to be the busiest*

```

# Here is a bar chart created for getting the most busiest weekday
present in our dataset.

weekday_names = ['Monday', 'Tuesday', 'Wednesday', 'Thursday',
'Friday', 'Saturday', 'Sunday']


# Grouping by pickup day and counting the number of pickups for each
day
pickup_counts =
data.groupby('pickup_day').size().reset_index(name='pickup_count')

pickup_counts['weekday_name'] = pickup_counts['pickup_day'].map(lambda
day: weekday_names[day % 7])

pickup_counts_by_weekday = pickup_counts.groupby('weekday_name')
['pickup_count'].sum().reset_index()

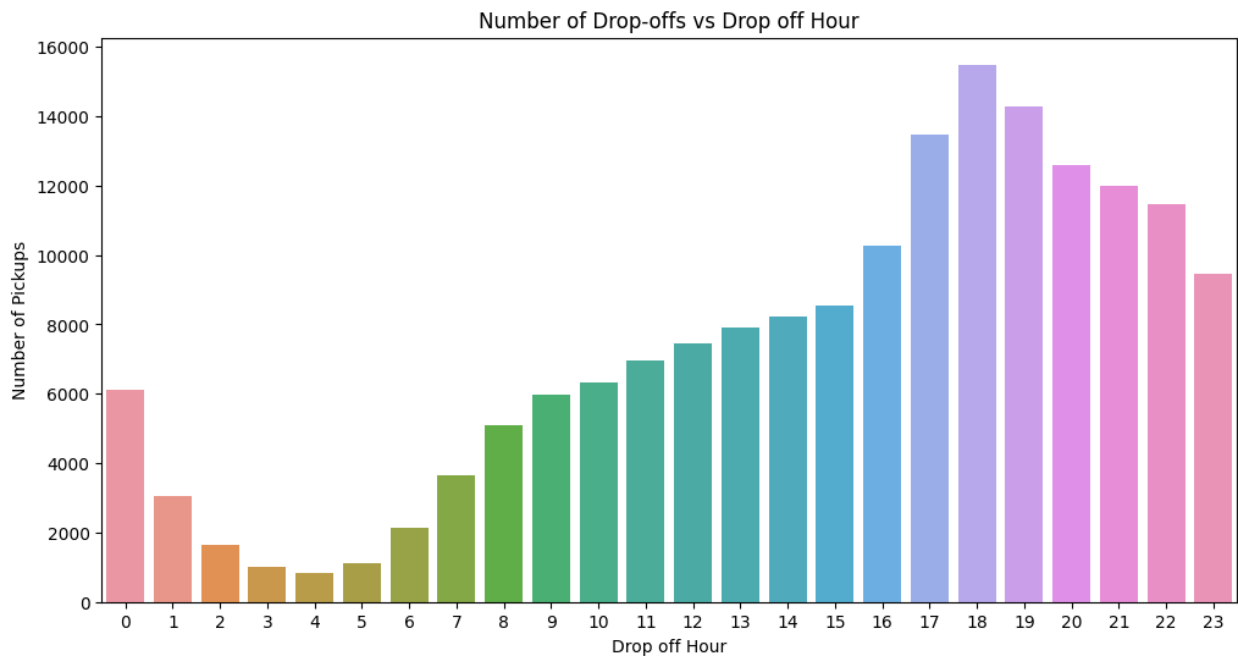
plt.figure(figsize=(10, 6))
plt.bar(pickup_counts_by_weekday['weekday_name'],
pickup_counts_by_weekday['pickup_count'], color='blue')
plt.xlabel('Day of the Week')
plt.ylabel('Number of Pickups')
plt.title('Number of Pickups per Day of the Week')
plt.show()

```



On our radar, Tuesday stands out as the busiest pickup day of the week, with Monday being the least active. A clear pattern emerges, highlighting the varying intensity of pickups throughout the week

```
# Here is a plot vs number of drop-offs vs drop off hour
plt.figure(figsize=(12, 6))
sns.countplot(x='dropoff_hour', data=data)
plt.title('Number of Drop-offs vs Drop off Hour')
plt.xlabel('Drop off Hour')
plt.ylabel('Number of Pickups')
plt.show()
```



When it comes to drop off hour, 2 am to 6 am in the morning is quiet idle where 5 pm to 11 pm in the evening seems to be the busiest

## Data Cleaning

data

	VendorID	passenger_count	trip_distance	RatecodeID	\
0	1	1.0	2.14	1.0	
1	0	1.0	2.70	1.0	
2	1	1.0	1.15	1.0	
3	0	1.0	0.40	1.0	
4	1	3.0	1.10	1.0	
...	...	...	...	...	
174995	1	3.0	3.45	1.0	
174996	1	1.0	9.44	1.0	
174997	0	1.0	2.40	1.0	

174998	1	1.0	4.71	1.0
174999	1	1.0	1.01	1.0

	store_and_fwd_flag	PULocationID	DOLocationID	payment_type
extra \				
0	N	120	9	Credit Card
2.5				
1	N	15	215	Credit Card
3.5				
2	N	167	223	Credit Card
0.0				
3	N	128	239	Credit Card
2.5				
4	N	203	52	Credit Card
1.0				
...	...	...	...	...
...				
174995	N	147	167	Credit Card
1.0				
174996	N	154	191	Cash
5.0				
174997	N	168	106	Credit Card
2.5				
174998	N	240	100	Credit Card
2.5				
174999	N	153	72	Credit Card
1.0				

	tip_amount	...	improvement_surcharge	total_amount	\
0	7.165589	...	1.0	20.64	
1	6.067401	...	1.0	25.55	
2	4.111547	...	1.0	17.64	
3	6.411079	...	1.0	12.80	
4	4.769377	...	1.0	18.00	
...	...	...	...	...	
174995	8.732495	...	1.0	28.08	
174996	0.283275	...	1.0	59.95	
174997	4.245354	...	1.0	33.50	
174998	10.479776	...	1.0	40.80	
174999	6.541699	...	1.0	16.32	

	congestion_surcharge	Airport_fee	pickup_day	pickup_month	\
0	2.5	0.00	28	6	
1	2.5	0.00	29	6	
2	2.5	0.00	30	6	
3	2.5	0.00	29	6	
4	2.5	0.00	29	6	
...	...	...	...	...	
174995	2.5	0.00	30	6	
174996	2.5	1.75	30	6	

174997	2.5	0.00	29	6
174998	2.5	0.00	29	6
174999	2.5	0.00	30	6

	pickup_hour	dropoff_day	dropoff_month	dropoff_hour
0	17	28	6	16
1	23	29	6	22
2	10	30	6	11
3	13	29	6	14
4	22	29	6	22
...	...	...	...	...
174995	22	30	6	22
174996	13	30	6	14
174997	11	29	6	12
174998	19	29	6	19
174999	21	30	6	22

[175000 rows x 21 columns]

*#importing the necessary libraries*

```
from sklearn.model_selection import train_test_split
from sklearn.compose import ColumnTransformer
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import OneHotEncoder
from sklearn.preprocessing import MinMaxScaler, StandardScaler
from sklearn.pipeline import Pipeline, make_pipeline
from sklearn.feature_selection import SelectKBest, chi2
from sklearn.tree import DecisionTreeClassifier
from sklearn.preprocessing import PolynomialFeatures
from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import mean_squared_error, r2_score
```

```
numerical_features = ['passenger_count', 'trip_distance', 'extra',
'tip_amount', 'tolls_amount',
'improvement_surcharge', 'congestion_surcharge', 'Airport_fee']
```

```
categorical_features = ['store_and_fwd_flag', 'payment_type']
```

```
numerical_features_most_frequent = ['RatecodeID', 'passenger_count']
```

```
#numerical_features_most_frequent2 = []
```

```
numerical_transformer = Pipeline(steps = [
    ('imputer', SimpleImputer(strategy="mean")),
    ('scaler', StandardScaler())
])
```

*# Pipeline for features with most frequent imputation strategy*

```
numerical_transformer_most_frequent = Pipeline(steps=[
    ('imputer_most_frequent',
SimpleImputer(strategy='most_frequent')),
    ('scaler', StandardScaler())
])
```

```

'''numerical_transformer_most_frequent2 = Pipeline(steps=[
    ('imputer_most_frequent2',
SimpleImputer(strategy='most_frequent')),
    ('scaler', StandardScaler()) # Assuming the same scaling for
these features
])'''

'''So, imputing missing values with the most frequent value and one-
hot encoding with handling of unknown categories.'''

categorical_transformer = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('onehot', OneHotEncoder(handle_unknown='ignore'))
])

preprocessor = ColumnTransformer(transformers=[
    ('num_most_frequent', numerical_transformer_most_frequent,
numerical_features_most_frequent),
    ('num', numerical_transformer, numerical_features),
    ('cat', categorical_transformer, categorical_features)
])

```

## SPLITTING THE DATASET INTO TRAIN AND TEST DATA

```

y = data['total_amount']
X_train,X_test, y_train , y_test =
train_test_split(data.drop(columns=['total_amount']),y,test_size
=0.2,random_state =42)

```

```
X_train.head()
```

	VendorID	passenger_count	trip_distance	RatecodeID	\
143961	1	1.0	7.79	1.0	
170292	1	1.0	0.79	1.0	
161029	1	1.0	0.29	2.0	
84006	0	1.0	0.60	1.0	
95628	0	1.0	1.90	1.0	

extra	store_and_fwd_flag	PULocationID	DOLocationID	payment_type
143961	N	181	174	Credit Card
1.0				
170292	N	250	226	Credit Card
2.5				
161029	N	236	251	unknown
0.0				
84006	N	83	166	Credit Card
3.5				
95628	N	70	35	Credit Card
5.0				

	tip_amount	tolls_amount	improvement_surcharge
congestion_surcharge \			
143961	7.956385	0.0	1.0
0.0			
170292	2.276785	0.0	1.0
2.5			
161029	1.062698	0.0	-1.0
-2.5			
84006	2.444217	0.0	1.0
2.5			
95628	5.163920	0.0	1.0
2.5			

	Airport_fee	pickup_day	pickup_month	pickup_hour
dropoff_day \				
143961	1.75	29	6	23
30				
170292	0.00	28	6	19
28				
161029	0.00	30	6	21
30				
84006	0.00	29	6	20
29				
95628	0.00	28	6	17
28				

	dropoff_month	dropoff_hour
143961	6	0
170292	6	18
161029	6	21
84006	6	21
95628	6	17

## MODEL BUILDING AND EVALUATION

```
#Applying Linear regression
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
model_lr = Pipeline(steps=[
    ('preprocessor_lr', preprocessor),
    ('model_lr', LinearRegression())
])

model_lr.fit(X_train, y_train)
y_pred_lr = model_lr.predict(X_test)

r2 = r2_score(y_test, y_pred_lr)
```

```

explained_variance_lr = r2 * 100
print("R2 score for linear regression is" ,explained_variance_lr)

R2 score for linear regression is 72.19555760853076

#Applying Polynomial Features
from sklearn.preprocessing import PolynomialFeatures
model_pf = Pipeline(steps=[
    ('preprocessor_pf', preprocessor),
    ('poly_features', PolynomialFeatures(degree=2)),
    ('model_pf', LinearRegression())
])

model_pf.fit(X_train, y_train)
y_pred_pf = model_pf.predict(X_test)
r2 = r2_score(y_test, y_pred_pf)
explained_variance_pf = r2 * 100
print("R2 score for polynomial regression is" ,explained_variance_pf)

R2 score for polynomial regression is 88.03033643454853

```

#### *Applying Hyperparameter Tuning*

```

from sklearn.model_selection import GridSearchCV

# Define the parameter grid for PolynomialFeatures
param_grid = {'poly_features__degree': [2,3]}

# Create the grid search
grid_search = GridSearchCV(model_pf, param_grid, cv=3,
    scoring='neg_mean_squared_error')

# Fit the grid search to your data
grid_search.fit(X_train, y_train)

# Get the best model from the grid search
best_model = grid_search.best_estimator_

# Print the best degree
best_degree = best_model.named_steps['poly_features'].degree
print("Best Polynomial Degree:", best_degree)

Best Polynomial Degree: 2

#Applying DecisionTree Regressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import PolynomialFeatures

model_dt = Pipeline(steps=[

```



```

        ('preprocessor_dt', preprocessor),
        ('model_dt', DecisionTreeRegressor())
    ])

```

```

model_dt.fit(X_train, y_train)
y_pred_dt = model_dt.predict(X_test)
r2 = r2_score(y_test, y_pred_dt)
explained_variance_dt = r2 * 100
print("R2 score for decision tree regression is", explained_variance_dt)

```

R2 score for decision tree regression is 91.7292913398568

#### *#Applying RandomForestRegressor*

```

from sklearn.ensemble import RandomForestRegressor
from sklearn.pipeline import Pipeline
from sklearn.metrics import r2_score

```

```

model_rf = Pipeline(steps=[
    ('preprocessor_rf', preprocessor),
    ('model_rf', RandomForestRegressor())
])

```

```

model_rf.fit(X_train, y_train)
y_pred_rf = model_rf.predict(X_test)
r2_rf = r2_score(y_test, y_pred_rf)
explained_variance_rf = r2_rf * 100
print("R2 score for Random Forest regression is", explained_variance_rf)

```

R2 score for Random Forest regression is 95.31685928889165

#### *#Applying XGBRegressor*

```

from xgboost import XGBRegressor
from sklearn.pipeline import Pipeline
from sklearn.metrics import r2_score

```

```

model_xgb = Pipeline(steps=[
    ('preprocessor_xgb', preprocessor),
    ('model_xgb', XGBRegressor())
])

```

```

model_xgb.fit(X_train, y_train)
y_pred_xgb = model_xgb.predict(X_test)
r2_xgb = r2_score(y_test, y_pred_xgb)
explained_variance_xgb = r2_xgb * 100
print("R2 score for XGBoost regression is", explained_variance_xgb)

```

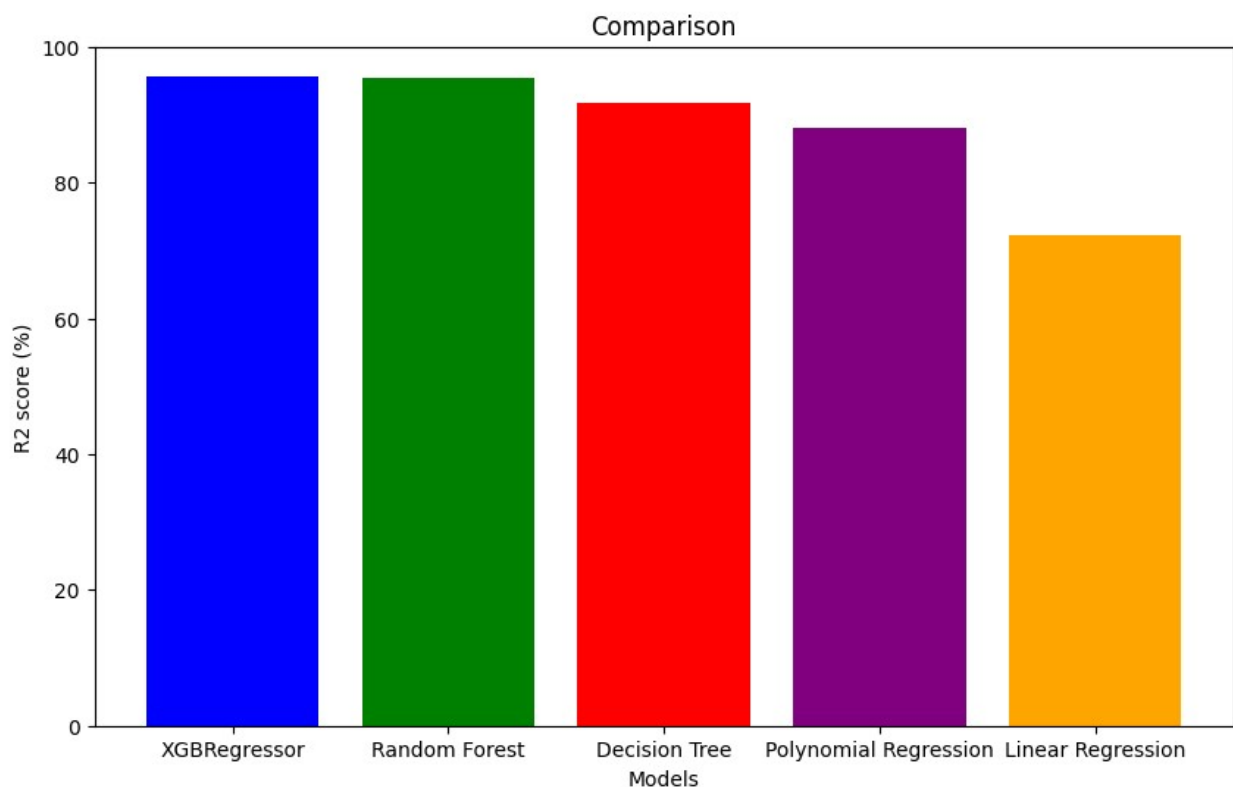
R2 score for XGBoost regression is 95.64346335641403

```

explained_variances = [explained_variance_xgb, explained_variance_rf,
explained_variance_dt, explained_variance_pf, explained_variance_lr]
model_names = ['XGBRegressor', 'Random Forest', 'Decision Tree',
'Polynomial Regression', 'Linear Regression']

plt.figure(figsize=(10, 6))
plt.bar(model_names, explained_variances, color=['blue', 'green',
'red', 'purple', 'orange'])
plt.title('Comparison')
plt.xlabel('Models')
plt.ylabel('R2 score (%)')
plt.ylim(0, 100)
plt.show()

```



**So , clearly XGBRegressor is our best model where Linear Regression being the worst one**

*Prediction*

```

y_pred_t=model_xgb.predict(data_t)
y_pred_t
array([34.858337, 24.224398, 14.648788, ..., 20.439754, 37.06559 ,
      16.91554 ], dtype=float32)

```

*Submission*

```
# Create a DataFrame with the 'ID' and 'total_amount' columns
submission_df = pd.DataFrame({
    'ID': range(1, len(y_pred_t) + 1), # Assuming index starts from 1
    'total_amount': y_pred_t
})

# Save the DataFrame to a CSV file
submission_df.to_csv('submission.csv', index=False)

submission_df.shape

(50000, 2)
```